

Computer Vision

Modelisation

Plan

Classification Results (CLS)



➤ CNN

- Convolution layer
- Pooling layer
- Fully-connected layer

➤ Transfert learning

- Purpose
- Strategies

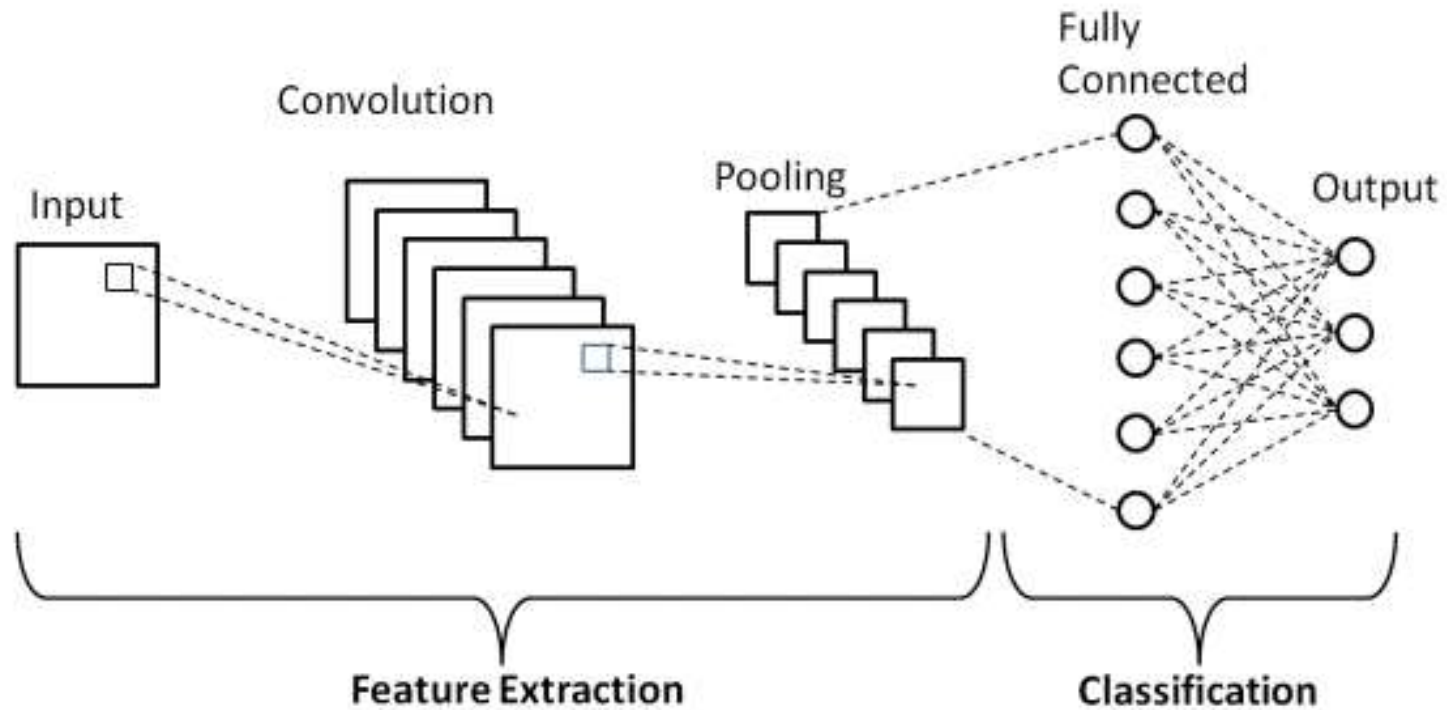
➤ Performance evaluation

- Hyperparameters
- Fitting



CNN

“Learn” the features

CNN



CNN: convolution layer



Activation function:

➤ ReLU: $f(x) = \text{Max}(0, x)$

CNN: convolution layer

Image (input):

- W : width (pixels)
- H : height (pixels)
- D : number of channels

Hyperparameters:

- K : number of filters
- F : filter size in pixels ($F \times F \times D$)
- S : sliding step (pixels)
- P : 0 padding (pixels)


Feature maps (output):

- $W_c : (W - F + 2P) / S + 1$
- $H_c : (H - F + 2P) / S + 1$
- $D_c : K$

Common choice:

- $F = 3$ or 5
- $S = 1$ or 2
- $P = 1$

CNN: pooling layer



Advantages:

- Dimension reduction
- Less sensitive to feature position

CNN: pooling layer

Feature maps (input):

- W : width (pixels)
- H : height (pixels)
- D : previously used number of filters

Hyperparameters:

- F : cell size ($F \times F$ pixels)
- S : pixels between cells

Feature maps (output):

- $W_p : (W - F) / S + 1$
- $H_p : (H - F) / S + 1$
- $D_p : D$


Common choice:

- $F = 2$ or 3
- $S = 2$

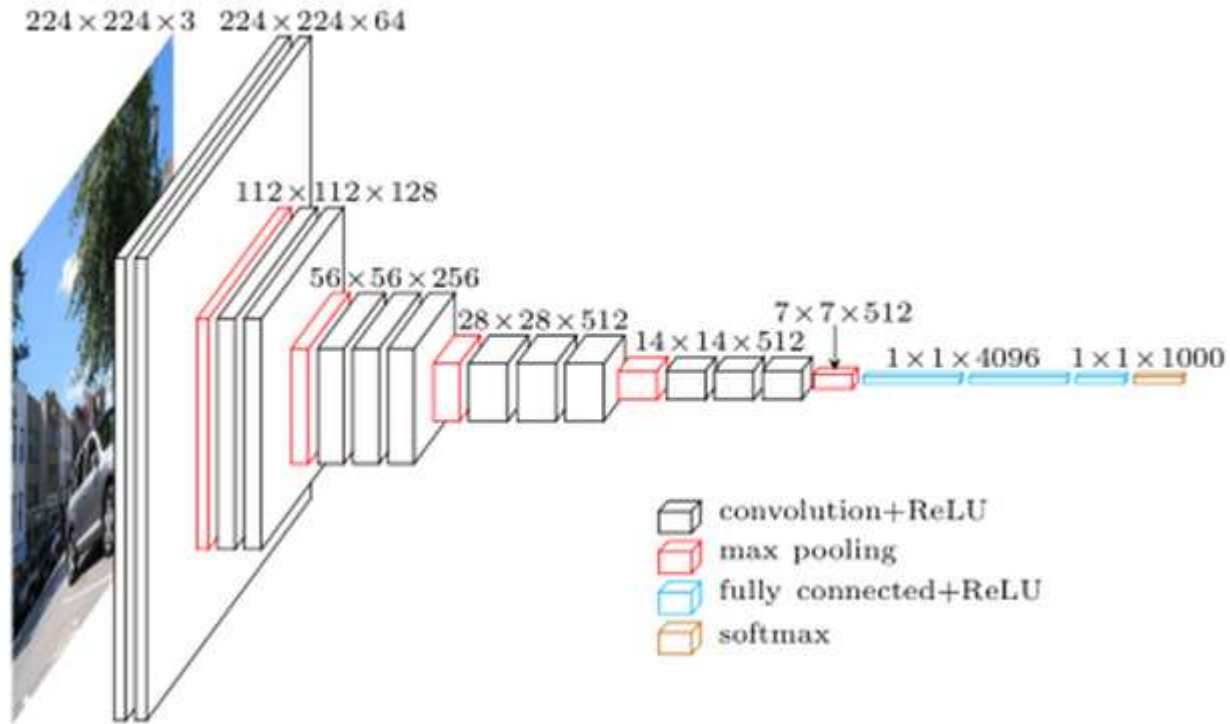
CNN: fully-connected layer

Options:

- Flatten: convert output to 1D vector
- Number of neurons
- Dropout: rate
- Activation:
 - ReLU
 - Classifier 2N: logistic regression
 - Classifier xN: Softmax



CNN: VGG16



CNN: VGG16

```
1 from keras.models import Sequential
2 from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
3
4 # Empty NN
5 my_VGG16 = Sequential()
6
7 #####
8 ### Convolution block ###
9 #####
10
11 # Add the first CNN layer with relu activation
12 my_VGG16.add(Conv2D(64, (3, 3), input_shape=(224, 224, 3), padding='same', activation='relu'))
13
14 # Add the second CNN layer with relu activation
15 my_VGG16.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
16
17 # Add the first pooling layer
18 my_VGG16.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
19
20 # Repeat the previous steps as often as needed...
21
22 #####
23 ### Fully-connected block ###
24 #####
25
26 # Convert 3D matrix to 1D vector
27 my_VGG16.add(Flatten())
28
29 # Add first fully-connected layer with relu activation
30 my_VGG16.add(Dense(4096, activation='relu'))
31
32 # Add second fully-connected layer with relu activation
33 my_VGG16.add(Dense(4096, activation='relu'))
34
35 # Add last fully-connected layer which is the classifier
36 my_VGG16.add(Dense(1000, activation='softmax'))
37
```

Transfer learning

Time to be lazy and... smart ! :)

Transfer learning: purpose

Weights:

- Convolution layer:
 - F : filter size
 - K : number of filters
 - $\text{weights} = F \times F \times K + K$
- Fully-connected:
 - W : width of input
 - H : height of input
 - K : number of filters
 - N : number of neurons
 - $\text{Weights} = W \times H \times K \times N + N$

**VGG16 = 138 357 544
weights !!!**

Transfer learning: strategies

The idea:

- Keep the architecture
- Replace the classifier

Strategies:

- Fine-tuning total:
 - Train all weights
 - Large dataset
- Feature extraction:
 - Fix all weights
 - Beside the ones of the classifier
 - Small / similar dataset
- Fine-tuning partial:
 - Fix some of the weights
 - Most likely in the first layers
 - Small / different dataset

Transfer learning: exemple

```
38
39 #####
40 ### Pre-trained exemple ###
41 #####
42
43 from keras.applications.vgg16 import VGG16
44
45 # Get VGG16 available in keras
46 model = VGG16()
47
48 from keras.preprocessing.image import load_img, img_to_array
49 from keras.applications.vgg16 import preprocess_input
50
51 # Load image and reshape image to match VGG16 dimensions expectation
52 img = load_img('cat.jpg', target_size=(224, 224))
53
54 # Convert to numpy array because keras processes images as such
55 img = img_to_array(img)
56
57 # A CNN expects (in general) a collection of images.
58 # Reshape the array to add the number of images
59 img = img.reshape((1, img.shape[0], img.shape[1], img.shape[2]))
60
61 # Preprocess images the same way the images used to train VGG16 were preprocessed.
62 img = preprocess_input(img)
63
64 # Predict the class (or label) of this image
65 y = model.predict(img)
66
```

Transfer learning: exemple

```
67
68 #####
69 ### Transfer learning ###
70 #####
71
72 from keras import Model
73
74 # Load the VGG16 trained on imagenet images without the fully-connected layer
75 model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
76
77 # Get the output of VGG16
78 x = model.output
79
80 # Add your new classifier
81 predictions = Dense(10, activation='softmax')(x)
82
83 # Design your new model
84 new_model = Model(inputs=model.input, outputs=predictions)
85
86 # Strat n°1:
87 for layer in new_model.layers:
88     layer.trainable = True
89
90 # Strat n°2:
91 for layer in new_model.layers:
92     layer.trainable = False
93
94 # Strat n°3: let's not train the first 5 layers (the largest)
95 for layer in new_model.layers[:5]:
96     layer.trainable = False
97
98 # Compile the new model
99 new_model.compile(loss='categorical_crossentropy', optimizer=optimizers.SGD(lr=0.0001, momentum=0.9), metrics=['accuracy'])
100
101 # Train it
102 model_info = new_model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=2)
103
104
```


Performance evaluation

The suspense is intense !

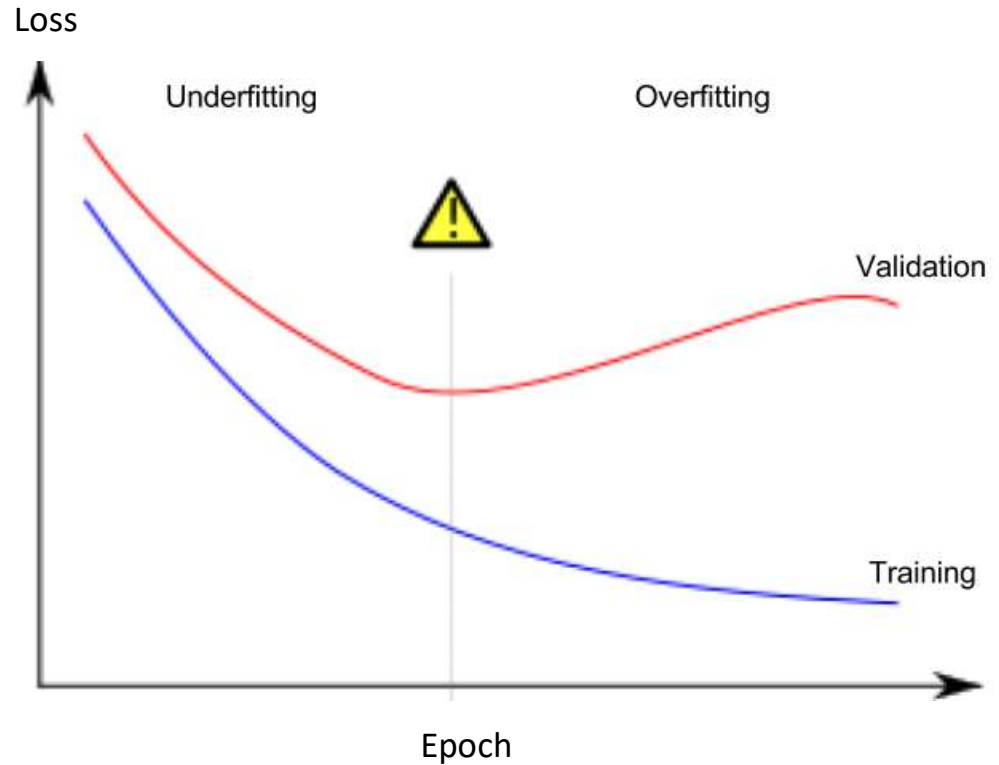
Hyperparameters

- Loss function
- Optimizer:
 - Learning rate
 - Momentum
- Metrics

```
57 #####
58 #####
59 ### Transfer learning ###
60 #####
61
62 from keras import Model
63
64 # Load the VGG16 trained on imagenet images without the fully-connected layer
65 model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
66
67 # Get the output of VGG16
68 x = model.output
69
70 # Add your new classifier
71 predictions = Dense(10, activation='softmax')(x)
72
73 # Design your new model
74 new_model = Model(inputs=model.input, outputs=predictions)
75
76 # Strat n°1:
77 for layer in new_model.layers:
78     layer.trainable = True
79
80 # Strat n°2:
81 for layer in new_model.layers:
82     layer.trainable = False
83
84 # Strat n°3: let's not train the first 5 layers (the largest)
85 for layer in new_model.layers[:5]:
86     layer.trainable = False
87
88 # Compile the new model
89 new_model.compile(loss="categorical_crossentropy", optimizer=optimizers.SGD(lr=0.0001, momentum=0.9), metrics=["accuracy"])
90
91 # Train it
92 model_info = new_model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=2)
93
94
```

Underfitting vs. Overfitting

```
my_model.fit(X_train, y_train,  
             early_stopping_rounds=5,  
             eval_set=[(X_valid, y_valid)],  
             verbose=False)
```



Practice

You know the drill

Ai design: dataset / concept

- **Select a dataset or an app idea**
- **Create your team !**
- **Define your strategy**
- **Prototype and test**

